

Course Introduction

Machine Learning

Daniele Loiacono



POLITECNICO
MILANO 1863

Information

□ Daniele Loiacono (Instructor)

- ▶ Contact: daniele.loiacono@polimi.it - +39 02 2399 3615
- ▶ Office: DEIB, room 150
- ▶ <https://webeep.polimi.it/course/view.php?id=8246>



□ Teaching Assistant: Alberto Maria Metelli



□ Exam

- ▶ Written test (closed-book)
 - Questions, exercises, code
 - See examples on WeBeep
 - **Late enrollment is NEVER accepted**
- ▶ Check the remote exam policy on WeBeep

Information (2)

- ❑ Weekly schedule
 - ▶ Tue, 14.15 – 16.15, Sala Conf. Emilio Gatti
 - ▶ Thu, 12.15 – 14.15, T23
- ❑ No streaming but lectures will be recorded
- ❑ Check [syllabus](#) for (tentative) info about topics of each lecture
- ❑ Practical classes
 - ▶ will cover exam-like exercises and practical examples
 - ▶ will present practical examples using **Python** (bring your laptop!)
- ❑ Interact
 - ▶ Feel free to ask questions
 - ▶ Use the forum of the course

References

- ❑ You will have access to all the materials used in classroom but slides are not an alternative to textbooks!
 - ▶ Most of the slides are inspired by the material of [prof. Restelli](#)

- ❑ Supervised Learning
 - ▶ Bishop, "Pattern Recognition and Machine Learning", Springer, 2006.
 - ▶ Hastie, Tibshirani, Friedman, "The Elements of Statistical Learning: Data Mining, Inference, and Prediction", Springer, 2009.
 - ▶ Mitchell, "Machine Learning", McGraw Hill, 1997.
 - ▶ Murphy, "Probabilistic Machine Learning: An Introduction", MIT Press, 2022

- ❑ Reinforcement Learning
 - ▶ Sutton and Barto, "Reinforcement Learning: an Introduction", MIT Press, 1998. New draft available at:
<http://www.incompleteideas.net/book/the-book-2nd.html>

What will you learn?

Goals

- ❑ Learn to correctly **model** machine learning problems
- ❑ Learn the **principles** and the **main techniques** of ML
- ❑ Learn how to **assess** the performances of ML models
- ❑ Learn **limitations** of ML techniques and how to **choose** the most appropriate one for your problem
- ❑ Provide the basic background to understand latest developments in this field

Topics

- ❑ Linear Regression
- ❑ Linear Classification
- ❑ Bias-Variance
- ❑ Model Selection
- ❑ PAC-Learning and VC dimension
- ❑ Kernel Methods
- ❑ Support Vector Machines
- ❑ Markov Decision Processes
- ❑ Dynamic Programming
- ❑ RL in finite MDPs
- ❑ Multi-armed bandit

What you will not learn?

Prerequisites

- ❑ Linear Algebra
 - ▶ Operations with matrix and vectors, eigenvalues, eigenvectors, etc.
- ❑ Probability and statistics
 - ▶ Distributions, confidence intervals, hypothesis test, bayesian statistics
- ❑ Optimization (basics)

- ❑ Basics understanding of Python (for practical classes)

- ❑ Where to find this?
 - ▶ Read Chapter 1-2 and Appendix B,C,E of textbook (Bishop, "Pattern Recognition and Machine Learning")
 - ▶ Check the recap lectures on Python, Linear Algebra and Probability
 - ▶ Chapter 1 of "Probabilistic Machine Learning: An Introduction"

Other courses

- ❑ A course of 5 credits is **not enough** to cover Machine Learning
- ❑ Fortunately, there are **other courses** that deal with other machine learning topics not covered in this course:
 - ▶ Data Mining and Text Mining
 - ▶ Soft Computing
 - ▶ Artificial Neural Networks and Deep Learning
 - ▶ Applied Statistics
 - ▶ Model Identification and Data Analysis

What is Machine Learning?

Why and when to apply it?

What is Machine Learning



“A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, improves with **experience E**” *Mitchell (1997)*

- ❑ ML is the sub-field of AI where the **knowledge** comes from:
 - ▶ **Experience**
 - ▶ **Induction**

- ❑ **However, Machine learning is not magic!**
 - ▶ You need to know **how it works**
 - ▶ You need to know **how to use it**
 - ▶ It can **extract** information from data, **not create** information

Why Machine Learning?


- ❑ We need computers to **make informed decisions** on new, unseen data
 - ▶ Often it is too difficult to design a set of meaningful rules
 - ▶ Machine learning allows to **automatically extract relevant information** from previous data and exploit it on new one
- ❑ Getting computers to **program themselves (automating automation)**
 - ▶ writing software is the **bottleneck**
 - ▶ let the **data** do the work instead



Machine Learning applications

- ❑ Machine learning is very popular today and has several applications:
 - ▶ Computer vision and robotics
 - ▶ Speech recognition
 - ▶ Biology and medicine
 - ▶ Finance
 - ▶ Information retrieval, Web search, ...
 - ▶ Entertainment and Videogames
 - ▶ Space exploration
 - ▶ Education
 - ▶ ...


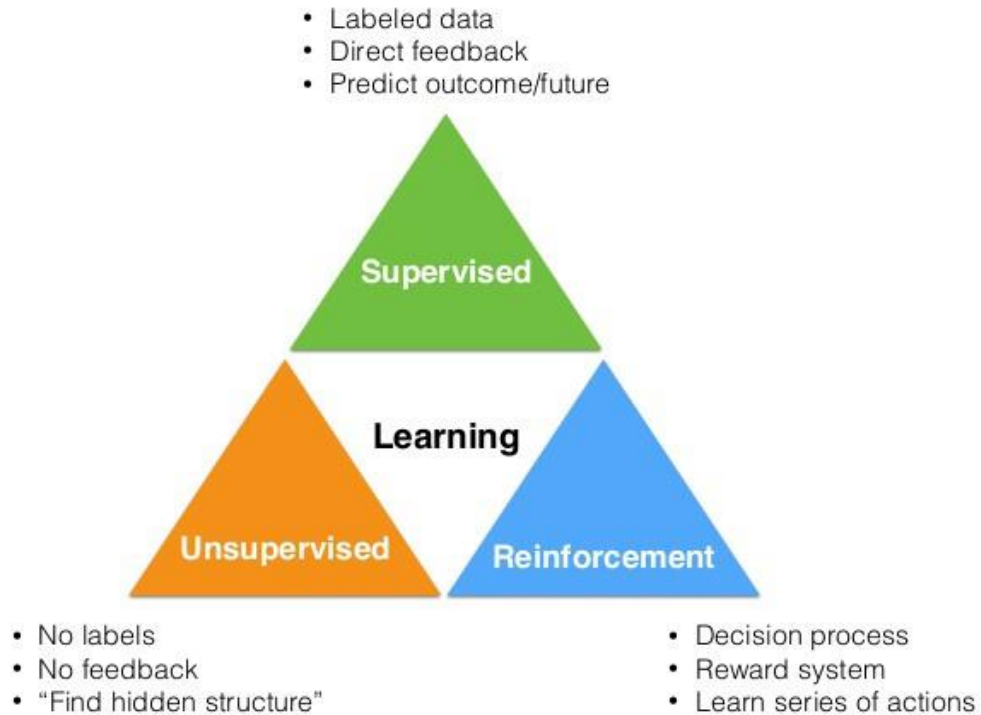
Machine Learning and other fields



Source: SAS, 2014 and PwC, 2016



Learning Paradigms in ML

□ Supervised Learning




Learning Paradigms in ML

□ Unsupervised Learning





Clustering



Dimensionality Reduction

Learning Paradigms in ML

□ Reinforcement Learning



Supervised Learning

□ Goal

- ▶ Learn from **data** a **model** that maps **known inputs** to **known outputs**
- ▶ Training set: $\mathcal{D} = \{\langle x, t \rangle\} \Rightarrow t = f(x)$

□ Tasks

- ▶ Classification
- ▶ Regression
- ▶ Probability estimation

□ Techniques


- ▶ Linear Models
- ▶ Artificial Neural Networks
- ▶ Support Vector Machines
- ▶ Decision trees
- ▶ etc.

An example of classification

Training Set





0



1



1



0


Testing



0



0




1



0

An example of regression


Training Set



78



34



25


Testing



78



81



41

Unsupervised Learning

□ Goal

- ▶ Learn **previously unknown patterns** and **efficient data representation**
- ▶ Training set: $\mathcal{D} = \{x\} \Rightarrow f(x)$


□ Tasks

- ▶ Dimensionality Reduction
- ▶ Clustering



□ Techniques

- ▶ K-means
- ▶ Self-organizing maps
- ▶ Principal Component Analysis
- ▶ etc.

An example of clustering



An example of dimensionality reduction



Reinforcement Learning

□ Goal

- ▶ Learning the **optimal policy**

- ▶ Training set: $\mathcal{D} = \{\langle x, u, x', r \rangle\} \Rightarrow \pi^*(x) = \arg \max_u \{Q^*(x, u)\}$

□ Problems

- ▶ Markov Decision Process (MDP)

- ▶ Partially Observable MDP (POMDP)

- ▶ Stochastic Games (SG)

to estimate

□ Techniques

- ▶ Q-learning

- ▶ SARSA

- ▶ Fitted Q-iteration

- ▶ etc.

An example of reinforcement learning

