

Model Evaluation, Selection and Ensembles

Machine Learning

Daniele Loiacono






Outline and References

□ Outline

- ▶ Bias-Variance Tradeoff (PRML 3.2, ESL 7.3)
- ▶ Model Assessment in Practice (ISL 5.1 and 6.1.3)
- ▶ Feature Selection (ESL 3.3)
- ▶ Dimensionality Reduction (PRML 12.1)
- ▶ Bagging and Boosting (PRML 14.2, 14.3)

□ References

- ▶ This slides are based on material of [prof. Marcello Restelli](#)
- ▶ [Pattern Recognition and Machine Learning, Bishop](#) [PRML]
- ▶ [Elements of Statistical Learning, Hastie et al.](#) [ESL]
- ▶ [Introduction to Statistical Learning, James et al.](#) [ISL]



Bias-Variance

How to evaluate a model?

- ❑ Given a dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$ and a model $\hat{t}_i = y(\mathbf{x}_i)$
- ❑ How do we choose a model y ?
- ❑ Is the loss function L computed on \mathcal{D} a good way to evaluate a model?
 - ▶ For example, for regression could we use $RSS_y = \sum_{\{\mathbf{x}_i, t_i\} \in \mathcal{D}} (t_i - y(\mathbf{x}_i))^2$
- ❑ No! This is the performance measure we want to compute:

$$\mathbb{E}[L] = \int \int L(t, y(\mathbf{x})) p(\mathbf{x}, t) d\mathbf{x} dt$$

- ▶ E.g., when using a squared loss function for regression:

$$\mathbb{E}[L] = \int \int (t - y(\mathbf{x}))^2 p(\mathbf{x}, t) d\mathbf{x} dt$$



We usually don't know $p(\mathbf{x}, t)$!
So what should we do?

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\mathbb{E}[(t - y(\mathbf{x}))^2] =$$

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\mathbb{E}[(t - y(\mathbf{x}))^2] = \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})]$$

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\begin{aligned}\mathbb{E}[(t - y(\mathbf{x}))^2] &= \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] + \mathbb{E}[y(\mathbf{x})^2] - \mathbb{E}[2ty(\mathbf{x})]\end{aligned}$$

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\begin{aligned}\mathbb{E}[(t - y(\mathbf{x}))^2] &= \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] + \mathbb{E}[y(\mathbf{x})^2] - \mathbb{E}[2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] \pm \mathbb{E}[t]^2 + \mathbb{E}[y(\mathbf{x})^2] \pm \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})]\end{aligned}$$

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\begin{aligned}\mathbb{E}[(t - y(\mathbf{x}))^2] &= \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] + \mathbb{E}[y(\mathbf{x})^2] - \mathbb{E}[2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] \pm \mathbb{E}[t]^2 + \mathbb{E}[y(\mathbf{x})^2] \pm \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \mathbb{E}[t]^2 + \text{Var}[y(\mathbf{x})] + \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})]\end{aligned}$$

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\begin{aligned}\mathbb{E}[(t - y(\mathbf{x}))^2] &= \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] + \mathbb{E}[y(\mathbf{x})^2] - \mathbb{E}[2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] \pm \mathbb{E}[t]^2 + \mathbb{E}[y(\mathbf{x})^2] \pm \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \mathbb{E}[t]^2 + \text{Var}[y(\mathbf{x})] + \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \text{Var}[y(\mathbf{x})] + (f(\mathbf{x}) - \mathbb{E}[y(\mathbf{x})])^2\end{aligned}$$

Bias-Variance Decomposition

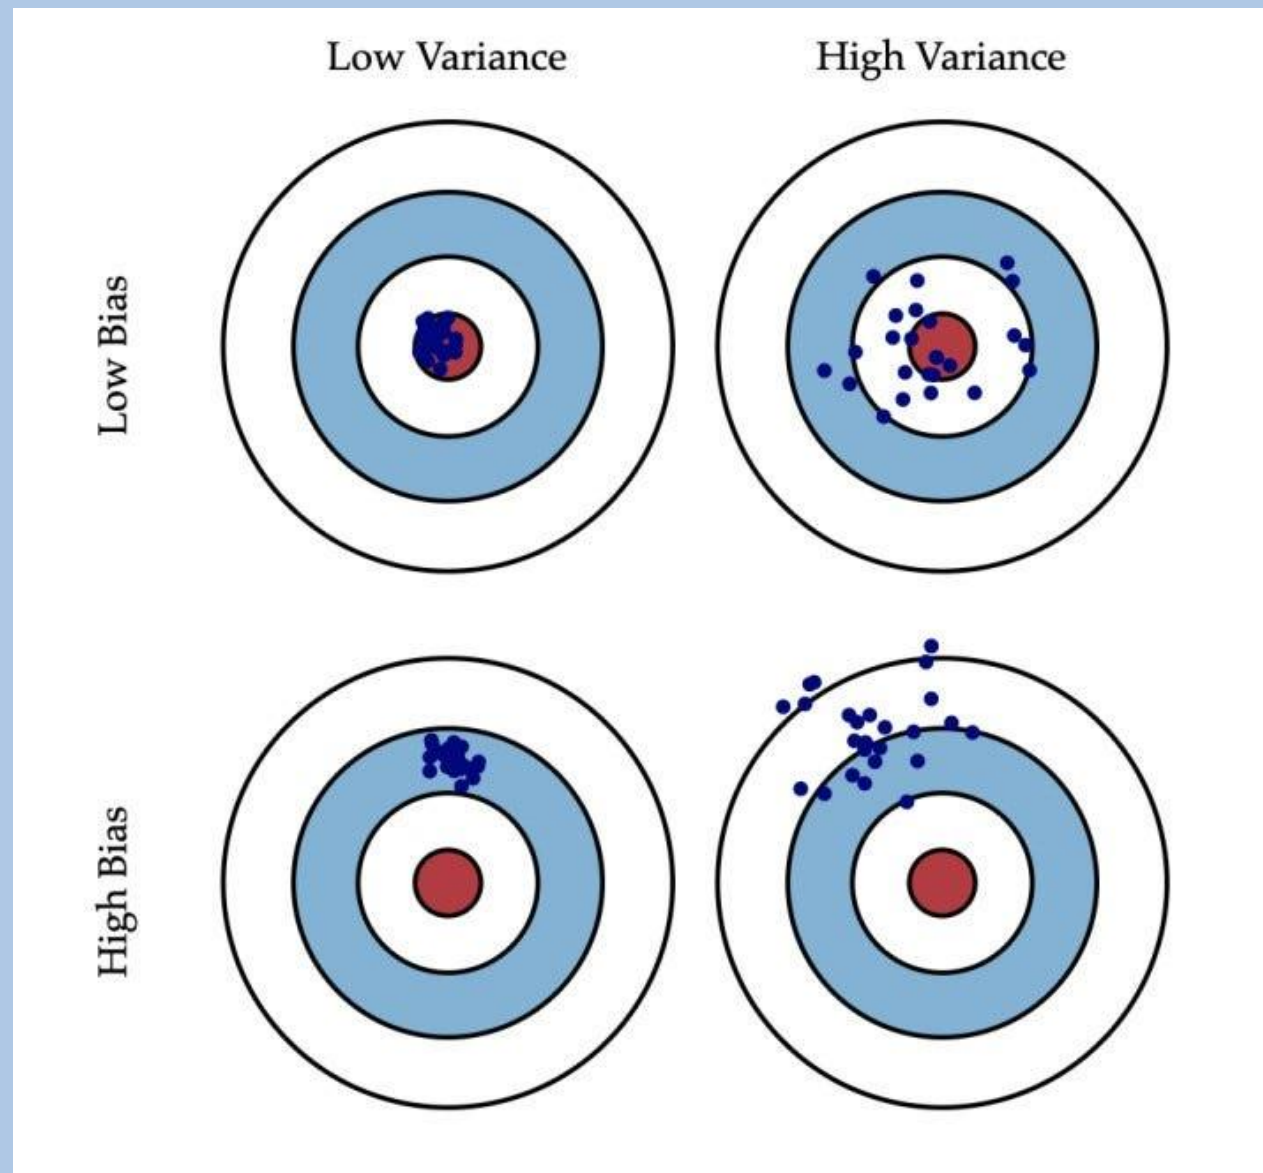
- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\begin{aligned}\mathbb{E}[(t - y(\mathbf{x}))^2] &= \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] + \mathbb{E}[y(\mathbf{x})^2] - \mathbb{E}[2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] \pm \mathbb{E}[t]^2 + \mathbb{E}[y(\mathbf{x})^2] \pm \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \mathbb{E}[t]^2 + \text{Var}[y(\mathbf{x})] + \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \text{Var}[y(\mathbf{x})] + (f(\mathbf{x}) - \mathbb{E}[y(\mathbf{x})])^2 \\ &= \text{Var}[t] + \text{Var}[y(\mathbf{x})] + \mathbb{E}[f(\mathbf{x}) - y(\mathbf{x})]^2\end{aligned}$$

Bias-Variance Decomposition

- The Bias-Variance is a framework to analyze the performance of models
- Definitions and assumptions
 - ▶ Data: $t_i = f(\mathbf{x}_i) + \varepsilon$, with $\mathbb{E}[\varepsilon] = 0$ and $\text{Var}[\varepsilon] = \sigma^2$
 - ▶ Model: $\hat{t}_i = y(\mathbf{x}_i)$ learned from a sampled dataset $\mathcal{D} = \{\mathbf{x}_i, t_i\}$
 - ▶ Performance: $\mathbb{E}[(t - y(\mathbf{x}))^2]$ (expected square error)
- Thus we can decompose the **expected square error** as:

$$\begin{aligned}\mathbb{E}[(t - y(\mathbf{x}))^2] &= \mathbb{E}[t^2 + y(\mathbf{x})^2 - 2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] + \mathbb{E}[y(\mathbf{x})^2] - \mathbb{E}[2ty(\mathbf{x})] \\ &= \mathbb{E}[t^2] \pm \mathbb{E}[t]^2 + \mathbb{E}[y(\mathbf{x})^2] \pm \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \mathbb{E}[t]^2 + \text{Var}[y(\mathbf{x})] + \mathbb{E}[y(\mathbf{x})]^2 - 2f(\mathbf{x})\mathbb{E}[y(\mathbf{x})] \\ &= \text{Var}[t] + \text{Var}[y(\mathbf{x})] + (f(\mathbf{x}) - \mathbb{E}[y(\mathbf{x})])^2 \\ &= \underbrace{\text{Var}[t]}_{\sigma^2} + \underbrace{\text{Var}[y(\mathbf{x})]}_{\text{Variance}} + \underbrace{\mathbb{E}[f(\mathbf{x}) - y(\mathbf{x})]^2}_{\text{Bias}^2}\end{aligned}$$



Bias-Variance Decomposition (2)

□ Model Variance


- ▶ If we sample several datasets \mathcal{D} we will learn different models $y(\mathbf{x})$
- ▶ variance measures the difference between each model learned from a particular dataset and what we expect to learn:

$$\text{variance} = \int \mathbb{E} \left[(y(\mathbf{x}) - \bar{y}(\mathbf{x}))^2 \right] p(\mathbf{x}) d\mathbf{x}$$

$\bar{y}(\mathbf{x}) = \mathbb{E}[y(\mathbf{x})]$

- ▶ Decreases with **simpler models**
- ▶ Decreases with **more samples**

Credit: <http://scott.fortmann-roe.com/docs/BiasVariance.html>



Bias-Variance Decomposition (3)

□ Model Bias


- bias measures the difference between **truth** (f) and what we expect to learn ($\mathbb{E}[y(\mathbf{x})]$):

$$\text{bias}^2 = \int (f(\mathbf{x}) - \bar{y}(\mathbf{x}))^2 p(\mathbf{x}) d\mathbf{x}$$

- Decreases with **more complex models**


- **Data noise** (σ^2) is the variance of data and does not depend neither from data sampling nor from the model complexity

Credit: <http://scott.fortmann-roe.com/docs/BiasVariance.html>




Bias-Variance Decomposition (4)

Credit: <http://scott.fortmann-roe.com/docs/BiasVariance.html>





Credit: <https://towardsdatascience.com/understanding-the-bias-variance-tradeoff-165e6942b229>



A case study: Bias-Variance for K-NN

□ Which is the best value for K?



Credit: <https://www.datacamp.com/community/tutorials/k-nearest-neighbor-classification-scikit-learn>

A case study: Bias-Variance for K-NN (2)

- Bias-Variance analysis allows to understand how K affects the performance:

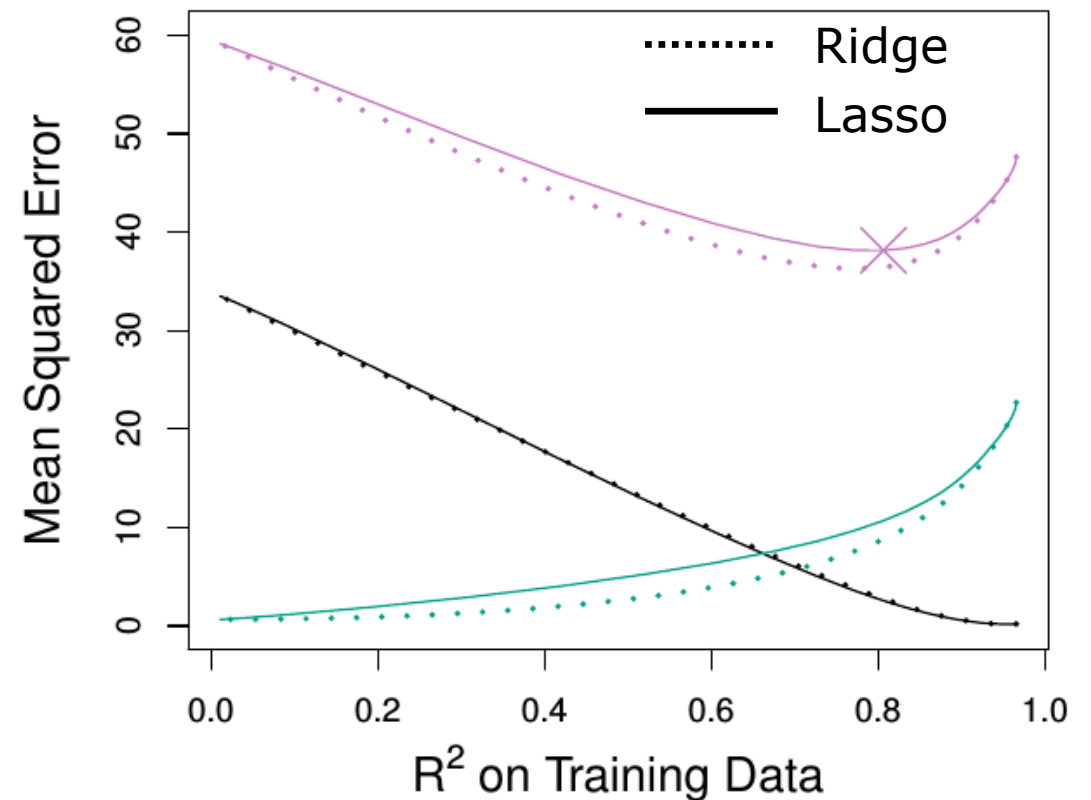
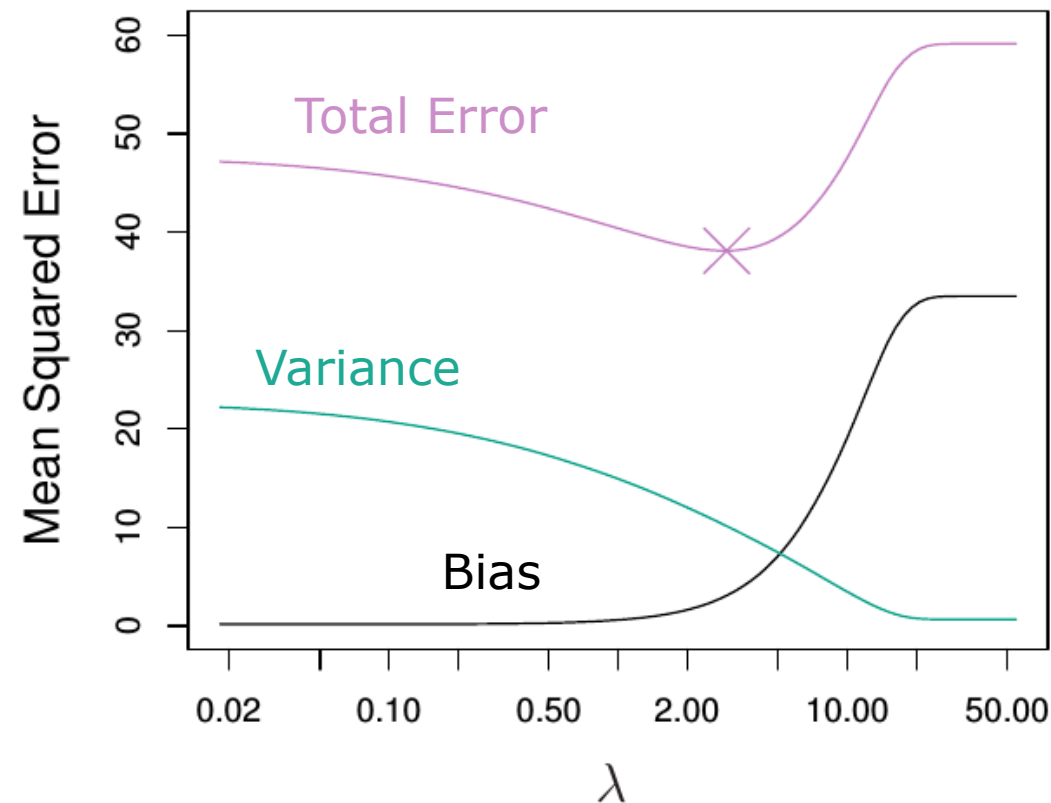
$$\mathbb{E}[(t^* - y(\mathbf{x}^*))^2] = \sigma^2 + \frac{\sigma^2}{K} + \left(f(\mathbf{x}^*) - \frac{1}{K} \sum_{i=1}^K f(\mathbf{x}_i) \right)^2$$

- ▶ The data noise σ^2 is the **irreducible error**
- ▶ The model **variance** $\frac{\sigma^2}{K}$ decreases as K increases
- ▶ The model **bias** (the last *squared* term) depends on **smoothness** of problem space, but in general increases as K increases (because it increases the distance of the neighbours used to compute y)

Regularization and Bias-Variance

- The Bias-Variance decomposition explains why regularization allows to improve the error on unseen data
- Lasso outperforms Ridge regression when few features are related to the output



45 of 45 features correlated to output



Regularization and Bias-Variance

- ❑ The Bias-Variance decomposition explains why regularization allows to improve the error on unseen data
- ❑ Lasso outperforms Ridge regression when few features are related to the output

2 of 45 features correlated to output



Model Assessment in Practice

Training Error


- Given $\mathcal{D} = \{\mathbf{x}_i, t_i\}$ with $i = 1, \dots, N$
- We can select a model based on the loss function L computed on \mathcal{D} :

- ▶ Regression

$$L_{train} = \frac{1}{N} \sum_{n=1}^N (t_n - y(\mathbf{x}_n))^2$$

- ▶ Classification

$$L_{train} = \frac{1}{N} \sum_{n=1}^N (I(t_n \neq y(\mathbf{x}_n)))$$



Prediction Error

- We already saw that **training error** does not provide a good estimate of the error made on new data, the **prediction error**


- ▶ Regression

$$L_{true} = \int \int (t - y(\mathbf{x}))^2 p(\mathbf{x}, t) d\mathbf{x} dt$$

- ▶ Classification

$$L_{true} = \int \int I(t \neq y(\mathbf{x})) p(\mathbf{x}, t) d\mathbf{x} dt$$

- Unfortunately, we often are not able to model $p(\mathbf{x}, t)$



So what to do in practice? Test error


- What can we do in practice?
 - ▶ Split randomly data into a **training set** and **test set**
 - ▶ Optimize model parameters using the **training set**
 - ▶ Estimate the prediction error using the **test set**

- Regression

$$L_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (t_n - y(\mathbf{x}_n))^2$$

- Classification

$$L_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (I(t_n \neq y(\mathbf{x}_n)))$$



So what to do in practice? Test error

□ What can we do in practice?

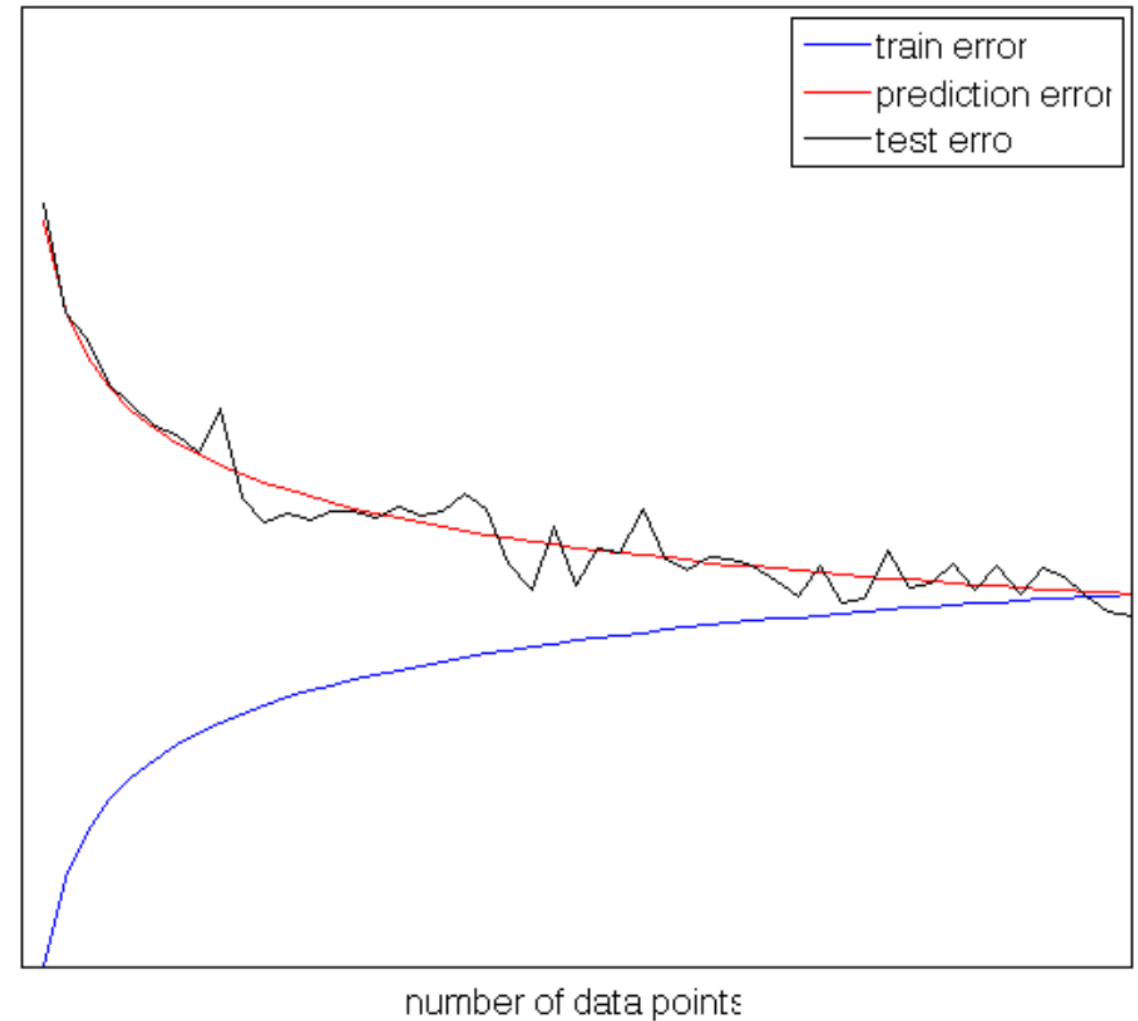
- ▶ Split randomly data into a **training set** and **test set**
- ▶ Optimize model parameters using the **training set**
- ▶ Estimate the prediction error using the **test set**

- Regression

$$L_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (t_n - y(\mathbf{x}_n))^2$$


- Classification

$$L_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (I(t_n \neq y(\mathbf{x}_n)))$$



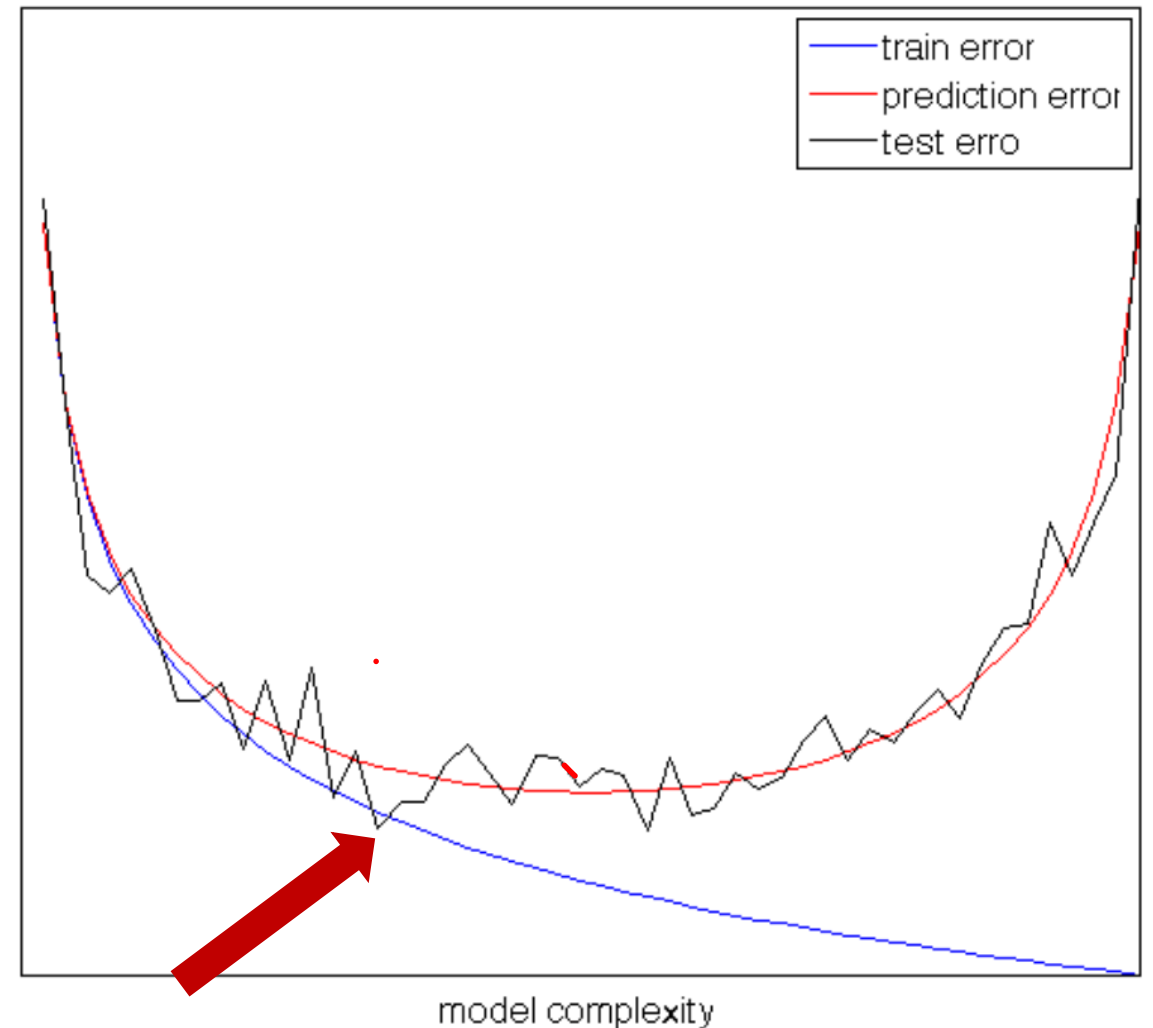
Analysis of train-test error

- The analysis of train-test errors allows to identify possible problems
 - ▶ **high bias**: training error is close to test error but they are both higher than expected
 - ▶ **high variance**: training error is smaller than expected and it slowly approaches the test error



Pitfalls of using test set

- ❑ Often data is limited and test error is usually small → prediction error can be either **overestimated** or **underestimated**
- ❑ Test error cannot be used for model selection → we can **overfit** test set
- ❑ Only if test set is **never used** for training and model selection, it can provide an **unbiased** estimate of prediction error
- ❑ So how to perform model selection?



Validation Set

- ❑ How to choose the best model and how do we set hyperparameters (e.g, α , λ)?
- ❑ We split data in three parts:
 - ▶ Training Data
 - ▶ Validation Data
 - ▶ Test Data
- ❑ How do we proceed?
 - ▶ We use **training data** to learn model parameters
 - ▶ For each model learned (i.e., different features and hyperparameters) we use **validation data** to compute the **validation error**
 - ▶ We select the model with the lowest **validation error** and finally use **test data** to estimate **prediction error**
- ❑ But...
 - ▶ to be reliable, validation data should be enough → less training data
 - ▶ we might **overfit validation data** and eventually not choose the best model

Leave-One-Out Cross Validation (LOO)

- For each sample $\{\mathbf{x}_i, t_i\} \in \mathcal{D}$
 - ▶ We train the model on $\mathcal{D} \setminus \{\mathbf{x}_i, t_i\}$
 - ▶ We compute the error of the resulting model on $\{\mathbf{x}_i, t_i\}$
- The estimate of the prediction error of our model will be the average of all the error computed using a single sample:

$$L_{LOO} = \frac{1}{N} \sum_{i=1}^N (t_i - y_{\mathcal{D}_i}(\mathbf{x}_i))^2$$

- ▶ Where $y_{\mathcal{D}_i}$ is the model trained on $\mathcal{D} \setminus \{\mathbf{x}_i, t_i\}$
- L_{LOO} provides an **almost unbiased** estimate of prediction error (**slightly pessimistic**)
- Unfortunately, LOO is **extremely expensive** to compute
 - ▶ As an example, even if training take just 1 second, computing LOO on 100K samples, would require 100K seconds (more than one day)!

K-Fold Cross Validation

- We randomly split the training data \mathcal{D} into k folds: $\mathcal{D}_1, \dots, \mathcal{D}_k$
- For each fold \mathcal{D}_i
 - ▶ We train the model on $\mathcal{D} - \{\mathcal{D}_i\}$
 - ▶ We compute the error on \mathcal{D}_i

$$L_{\mathcal{D}_i} = \frac{k}{N} \sum_{(\mathbf{x}_n, t_n) \in \mathcal{D}_i} (t_n - y_{\mathcal{D} \setminus \mathcal{D}_i}(\mathbf{x}_n))^2$$

- Finally, we estimate the prediction error as the average error computed:

$$L_{k-fold} = \frac{1}{k} \sum_{i=1}^k L_{\mathcal{D}_i}$$

- L_{k-fold} provides a **slightly biased** estimate of prediction error (**pessimistic**) but it is much cheaper to compute (usually $k=10$ is used)

Complexity-Adjusted Model Evaluation

- ❑ Other metrics are used to evaluate models adjusting their training error based on their complexity:

- ▶ Mallows's C_p : $C_p = \frac{1}{N} (RSS + 2M\hat{\sigma}^2)$
- ▶ Akaike Information Criteria: $AIC = -2\ln L + 2M$
- ▶ Bayesian Information Criteria: $BIC = -2\ln L + M\ln(N)$
- ▶ Adjusted R^2 : $AdjustedR^2 = 1 - \frac{RSS/(N-M-1)}{TSS/(N-1)}$

M: number of parameters; N: number of samples; L: loss function;
 $\hat{\sigma}^2$: estimate of noise variance;
RSS: residual sum of squares; TSS: total sum of squares

- ❑ AIC and BIC are generally used when maximizing the log-likelihood
- ❑ BIC generally penalize more than AIC model complexity


How to choose the right model complexity?

Curse of dimensionality


- ❑ Adding a feature means an exponential increase of volume of the input space
- ❑ Challenges
 - ▶ Computational cost
 - ▶ Amount of data...
 - ▶ ... large model variance (overfitting)

Curse of dimensionality


- Adding a feature
- Challenges
 - ▶ Computational cost
 - ▶ Amount of data
 - ▶ ... large number of features




1D



2D



3D



4D

input space

Curse of dimensionality

- ❑ Adding a feature means an exponential increase in volume of the input space
- ❑ Challenges
 - ▶ Computational cost
 - ▶ Amount of data...
 - ▶ ... large model variance (overfitting)
- ❑ A common pitfall is to think that using more features is always better:
 - ▶ Is $y = w_0 + w_1x + w_2x^2$ always better than $y = w_0 + w_1x$, isn't it?
 - ▶ In fact, we can always set $w_2 = 0$
 - ▶ No! Increasing the number of features increases the probability of overfitting the data (also because I could not have enough data for a larger feature space)!


Reducing the variance

- ❑ We want to select the model with the lowest prediction error
- ❑ This can be achieved by reducing the variance of the model:
 - ▶ **Feature Selection**: we should design the feature space by selecting the **most effective subset** of all the possible features
 - ▶ **Dimensionality Reduction**: the input space can be mapped to a **lower-dimensional space**
 - ▶ **Regularization**: the values of the parameters are **shrunk toward zero**
- ❑ These three approaches are not necessarily mutually exclusive and they can be used together

Feature Selection

Best Subset Selection: a brute force

- ❑ The simplest idea seems to compare all the possible combinations of the features
- ❑ Assuming we have M features, for each $k = 1, \dots, M$:
 - ▶ Train all the $\binom{M}{k} = \frac{M!}{k!(M-k)!}$ models with exactly k features
 - ▶ Chose the best model
- ❑ Computation unfeasible!



Feature Selection in Practice



- ❑ **Filter:** features are ranked **independently** based on some evaluation metrics (e.g., correlation, variance, information gain, etc.) and the top k are selected
 - ▶ Very fast but fails to capture any subset of **mutually dependent features**
- ❑ **Embedded:** feature selection is performed as a step of the machine learning approach used (e.g., lasso, decision trees, etc.)
 - ▶ Not expensive but the features identified are **specific** to the learning approach
- ❑ **Wrapper:** a search algorithm is used to find a subset of features that are evaluated by training a model with them and assessing its performance
 - ▶ Either a simpler model or a simple machine learning approach can be used to evaluate the features
 - ▶ Greedy algorithms are generally used to search the best subset of features

Dimensionality Reduction

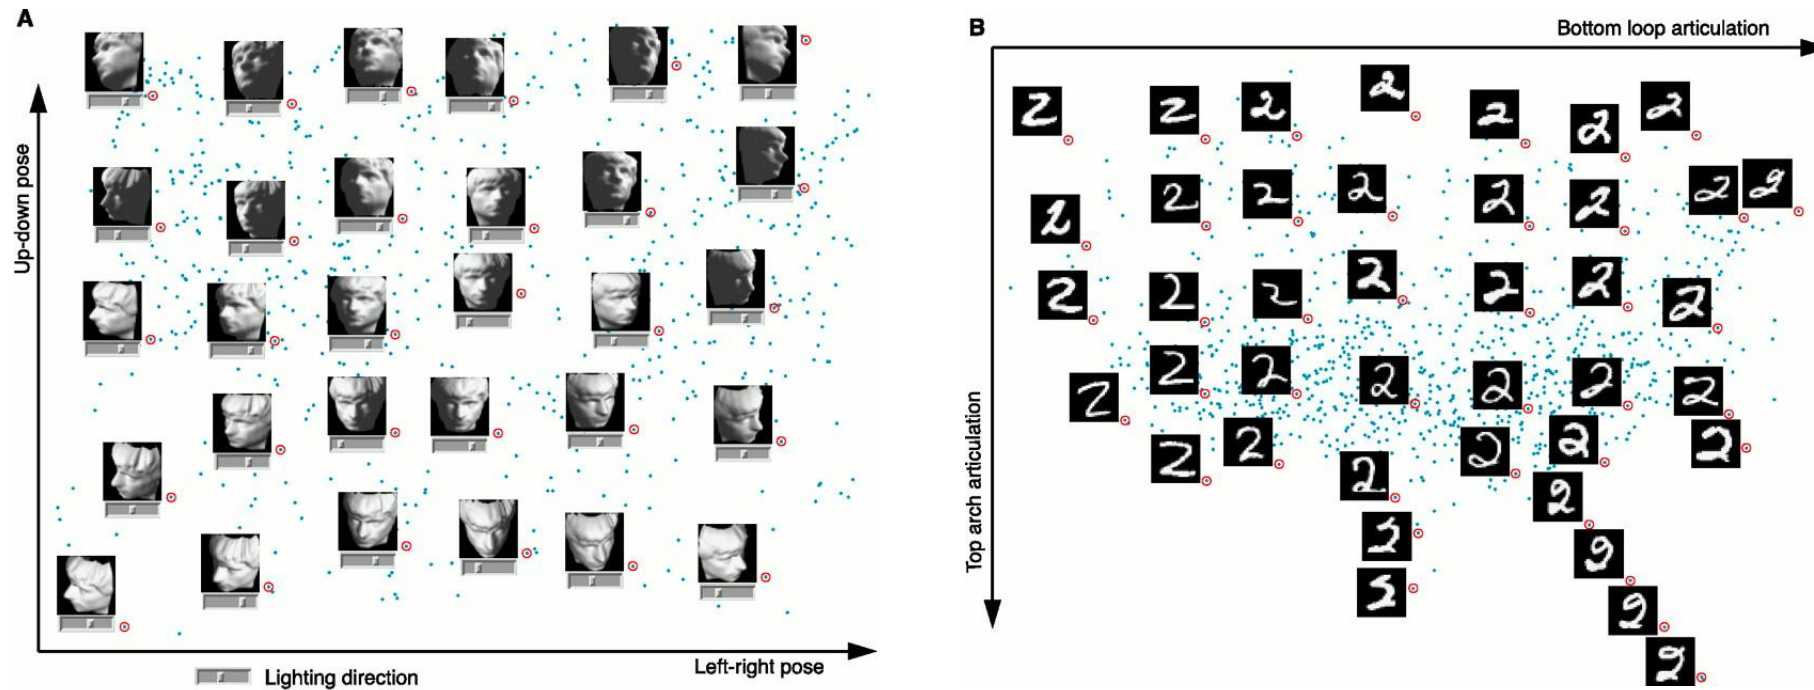
Dimensionality Reduction

- ❑ Dimensionality reduction aims at reducing the dimensions of input space, but it differs from feature selection in two major respects:
 - ▶ it uses **all** the features and **maps** them into a lower-dimensionality space
 - ▶ it is an **unsupervised** approach
- ❑ There are many methods to perform dimensionality reduction:
 - ▶ **Principal Component Analysis (PCA)**
 - ▶ **Indipendent Component Analysis (ICA)**
 - ▶ **Self-Organizing Maps**
 - ▶ **Autoencoders**
 - ▶ **ISOMAP**
 - ▶ **t-SNE**
 - ▶ ...

PCA underlying idea



DR: an example with complex data



*[A global geometric framework for nonlinear dimensionality reduction.
Tenenbaum, de Silva, and Langford. Science, 290(5500):2319–2323, 2000.]*

Bagging and Boosting

Improving the Bias-Variance Tradeoff

- ❑ So far we saw how to reduce the variance, searching the best tradeoff with the increased bias...
- ❑ ... but it is possible to reduce the variance **without increasing bias**?
- ❑ Or is it possible to **reduce the bias**?
- ❑ We can achieve these results by using two **ensemble methods** that consist of learning **several models** and **combine** them:
 - ▶ Bagging
 - ▶ Boosting

Which is the idea behind bagging?

- ❑ Let assume to have N datasets and to learn from them N models, y_1, y_2, \dots, y_N
- ❑ Now let us compute an aggregate model as $y_{AGG} = \frac{1}{N} \sum_{i=1}^N y_i$
- ❑ If the datasets are **independent**, the model variance of y_{AGG} will be $1/N$ of the model variance of the single model y_i
 - ▶ To have an intuition of this, let assume to have random variable $\bar{x} = \frac{1}{N} \sum_N x$

$$Var(\bar{x}) = \frac{1}{N^2} \sum_N Var(x) = \frac{1}{N} Var(x)$$

- ❑ However, we generally do not have N datasets!
- ❑ So, what can we do?


Bagging

- ❑ Bagging stands for Bootstrap Aggregation:
 - ▶ Generate N datasets applying **random sampling with replacement**
 - ▶ Train a model (classification or regression model) from **each** dataset generated
 - ▶ To compute the **prediction** for new samples, apply all the trained models and combine the outputs with **majority voting** (classification) or **averaging** (regression)
- ❑ Bagging is generally helpful and reduce the variance, although the sampled datasets are **not independent**
 - ▶ It helps with **unstable learners**, i.e., learners that change significantly with even small changes in the dataset (**low bias and high variance**)
 - ▶ It helps when we have a lot of overfitting (**low bias and high variance**)
 - ▶ It does not help when learners is **robust**, i.e., not sensitive to change in data (**usually higher bias but lower variance**)

What is Boosting?

- ❑ The goal of boosting is to achieve a **small bias** by using on simple (**weak**) learners
- ❑ The key idea behind boosting is to **iteratively train** a series of **weak learners**, with each iteration focusing on the samples that were misclassified in the previous iteration.
- ❑ As final result, an **ensemble model** is built by **combining** the outputs of all the weak learners trained

AdaBoost



Bagging vs Boosting

Bagging

- ❑ Reduces variance
- ❑ Not good for stable learners
- ❑ Can be applied with noisy data
- ❑ Usually helps but the difference might be small
- ❑ Naturally parallel

Boosting

- ❑ Reduces bias (generally without overfitting)
- ❑ Works with stable learners
- ❑ Might have problem with noisy data
- ❑ Not always helps but it can make the difference
- ❑ Serial