

Temporal-Difference Learning

Machine Learning

Daniele Loiacono



POLITECNICO
MILANO 1863

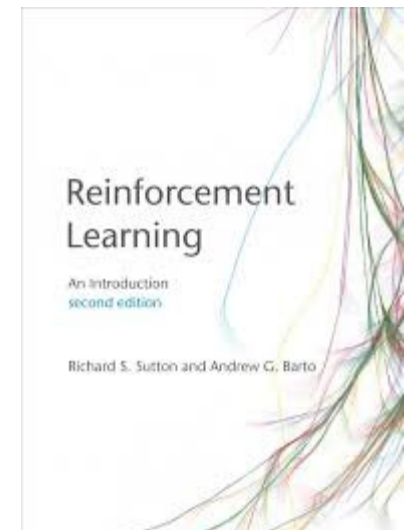
Outline and References

□ Outline

- ▶ TD(0)
- ▶ SARSA
- ▶ Q-Learning


□ References

- ▶ [Reinforcement Learning: An Introduction](#) [RL Chapter 6 and 7]
- ▶ [Sample-based Learning Methods](#) (Coursera)



Why temporal-difference?

- ❑ DP requires the knowledge of the MDP dynamics
- ❑ MC learns from experience but requires complete episodes to perform the updates
 - ▶ It can be applied only to episodic task
 - ▶ Even in episodic task it might have problems:



TD(0)

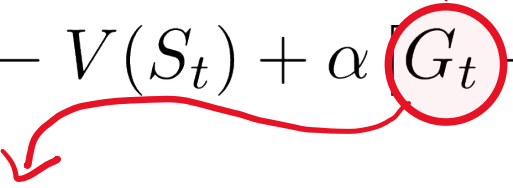
Temporal-Difference Policy Evaluation – TD(0)

- Temporal-Difference combines MC (**model-free**) with DP (**bootstrapping**):

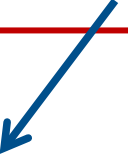
$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

Temporal-Difference Policy Evaluation – TD(0)

- Temporal-Difference combines MC (model-free) with DP (bootstrapping):

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$
$$G_t \approx R_{t+1} + \gamma V(S_{t+1})$$




$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$


Temporal-Difference Error (δ_t)

TD(0) Policy Evaluation

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal

TD(0) Policy Evaluation

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

Initialize S

Loop for each step of episode.

$A \leftarrow$ action given by π for S

Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

until S is terminal

S,A,R,S,A,S,A,R,S,A,S,A,R,...

TD(0) Policy Evaluation

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

Initialize S

Loop for each step of episode.

$A \leftarrow$ action given by π for S

Take action A , observe R, S'


$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

until S is terminal




update
 $S, A, R, S, A, S, A, R, S, A, S, A, R, \dots$

Random Walk Example



- ❑ Let C be the starting state
- ❑ Let π be left or right with equal probability in all states (random policy)
- ❑ Reward +1 on **right** termination, 0 otherwise
- ❑ Assuming $\gamma = 1$, $V_{\pi}(s)$ represents the probability of ending on the right side

Random Walk Example – Value Function



MC vs TD

- ❑ TD can learn **before** knowing the final outcome
 - ▶ TD can learn **after every step**
 - ▶ MC must wait **until end of episode** before return is known
- ❑ TD can learn **without** the final outcome
 - ▶ TD can learn from **incomplete sequences**
 - ▶ MC can only learn from **complete sequences**
- ❑ TD works in continuing (non-terminating) environments MC only works for episodic (terminating) environments

MC vs TD: Bias-Variance

❑ MC target has lower bias

- ▶ MC return $R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_{t+T}$ is an **unbiased** estimate of $V_{\pi}(S_t)$
- ▶ TD target $R_{t+1} + \gamma V(S_{t+1})$ is a **biased** estimate of $V_{\pi}(S_t)$, as $V(S_{t+1}) \neq V_{\pi}(S_{t+1})$

❑ TD target has lower variance

- ▶ Return depends on many random actions, transitions, rewards
- ▶ TD target depends on one random action, transition, reward

❑ Overall

- ▶ MC works well with function approximation and it is not very sensitive to initial values
- ▶ TD has problem with function approximation and it is more sensitive to initial values

Sampling and Bootstrapping: overview

- ❑ **Bootstrapping:** update involves an estimate
 - ▶ MC does not bootstrap
 - ▶ DP bootstraps
 - ▶ TD bootstraps
- ❑ **Sampling:** update does not involve an expected value
 - ▶ MC samples
 - ▶ DP does not sample
 - ▶ TD samples

SARSA

Let's go back to MC Policy Iteration


□ Evaluation

$$Q_{\pi}(s, a) \approx \text{average}[G_t | S_t = s, A_t = a]$$

$$G_t \approx R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$$

□ Improvement


$$\pi'(s) = \arg \max_a Q_{\pi}(s, a)$$



On-Policy Control with SARSA

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

- Evaluation
 - ▶ SARSA
- Improvement
 - ▶ ϵ -Greedy Policy Improvement



On-Policy Control with SARSA

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

On-Policy Control with SARSA

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ε -greedy)

Loop for each step of episode:

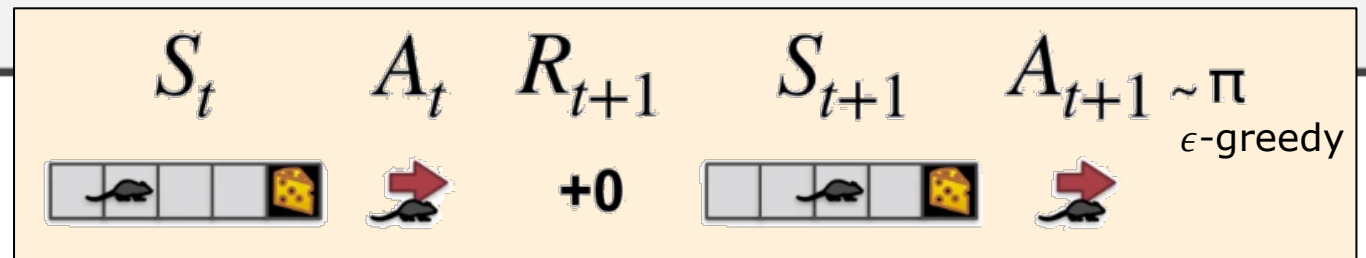
Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ε -greedy)


$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal



Example: SARSA on Windy Gridworld



Q-Learning

Q-Learning: Off-Policy TD Control

- SARSA, as Policy Iteration in DP, is based on Bellman Expectation Equation

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

$$Q_{\pi}(s, a) = \sum_{s', r} p(s', r | s, a) \left(r + \gamma \sum_{a'} \pi(a' | s') Q_{\pi}(s', a') \right)$$

- Q-Learning, as Value Iteration in DP, is based on Bellman Optimality Equation

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right)$$

$$Q^*(s, a) = \sum_{s', r} p(s', r | s, a) \left(r + \gamma \max_{a'} Q^*(s', a') \right)$$

Q-Learning: Off-Policy TD Control

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)



 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

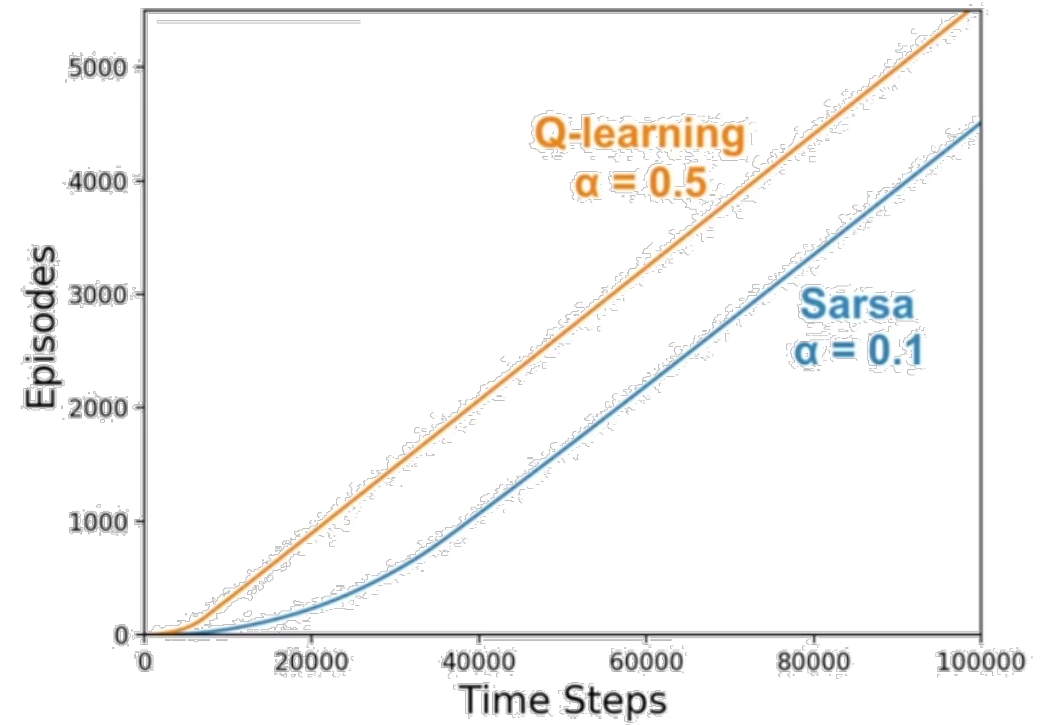

$S \leftarrow S'$

 until S is terminal

Q-Learning vs SARSA: Windy Gridworld



Q-Learning vs SARSA: Windy Gridworld



Q-Learning vs SARSA: Cliff Walking Environment

